

Assignment 2 - dplyr exploration

Lulu Pei

22/09/2019

Overview

Before playing around with the various `dplyr` commands, we must define a dataset. In this case, we will be looking at the `gapminder` dataset. To access the `gapminder` dataset, we need to load the `gapminder` package. To be able to call and use `dplyr` commands, we must load the `tidyverse` package as well.

```
library(gapminder)
library(tidyverse)
```

Question 1 - Basic dplyr

1.1 - Filtering

Filter `gapminder` data to contain only observations from Canada, the United States, and Mexico in the 1970s.

```
gapminder %>%
  filter(country == 'Canada' | country == 'United States' | country == 'Mexico',
         year >= 1970 & year < 1980) %>%
  knitr::kable()
```

country	continent	year	lifeExp	pop	gdpPercap
Canada	Americas	1972	72.880	22284500	18970.571
Canada	Americas	1977	74.210	23796400	22090.883
Mexico	Americas	1972	62.361	55984294	6809.407
Mexico	Americas	1977	65.032	63759976	7674.929
United States	Americas	1972	71.340	209896000	21806.036
United States	Americas	1977	73.380	220239000	24072.632

1.2 - Selecting

Let's select only `country` and `gdpPercap` variables from our filtered subset of the `gapminder` dataset.

```
gapminder %>%
  filter(country == 'Canada' | country == 'United States' | country == 'Mexico',
         year >= 1970 & year < 1980) %>%
  select(country, gdpPercap) %>%
  knitr::kable()
```

country	gdpPercap
Canada	18970.571
Canada	22090.883
Mexico	6809.407
Mexico	7674.929
United States	21806.036
United States	24072.632

1.3 - Mutating

Suppose we want to look at all the countries that have ever experienced a drop in life expectancy between 1952 and 2007. Let's define a new variable `lifeExp_change`, equaling the difference between life expectancy at one time point and life expectancy at the time point before.

```
gapminder %>%
  group_by(country) %>%
  mutate(lifeExp_change = lifeExp - lag(lifeExp, order_by = year)) %>%
  filter(lifeExp_change < 0)
```

```
# A tibble: 102 x 7
# Groups:   country [52]
  country continent year lifeExp      pop gdpPercap lifeExp_change
  <fct>      <fct>   <int>   <dbl>   <int>    <dbl>         <dbl>
1 Albania Europe    1992    71.6  3326498   2497.         -0.419
2 Angola  Africa    1987    39.9  7874230   2430.         -0.036
3 Benin   Africa    2002    54.4  7026113   1373.         -0.371
4 Botswana Africa    1992    62.7  1342614   7954.         -0.877
5 Botswana Africa    1997    52.6  1536536   8647.        -10.2
6 Botswana Africa    2002    46.6  1630347  11004.         -5.92
7 Bulgaria Europe    1977    70.8  8797022   7612.         -0.09
8 Bulgaria Europe    1992    71.2  8658506   6303.         -0.15
9 Bulgaria Europe    1997    70.3  8066057   5970.         -0.87
10 Burundi Africa    1992    44.7  5809236    632.         -3.48
# ... with 92 more rows
```

We can see that many countries have experienced a drop in life expectancy at some point between 1952 and 2007; however, what if we are only interested in countries that have experienced an overall life expectancy drop between the most recent year (2007) and the earliest year (1952).

```
gapminder %>%
  filter(year == 1952 | year == 2007) %>%
  group_by(country) %>%
  mutate(lifeExp_change = lifeExp - lag(lifeExp)) %>%
  filter(lifeExp_change < 0) %>%
  knitr::kable()
```

country	continent	year	lifeExp	pop	gdpPercap	lifeExp_change
Swaziland	Africa	2007	39.613	1133066	4513.4806	-1.794
Zimbabwe	Africa	2007	43.487	12311143	469.7093	-4.964

From this output, we can see that the only countries that experienced an overall decline in life expectancy were Swaziland and Zimbabwe, with a decline of 1.794 and 4.964 years, respectively.

1.4 - Slicing

Now, let's filter the `gapminder` dataset to show the maximum GDP per capita experienced by each country.

```
gapminder %>%
  select(country, gdpPercap) %>%
  group_by(country) %>%
  slice(which.max(gdpPercap)) %>%
  rename(max_gdpPercap = gdpPercap)
```

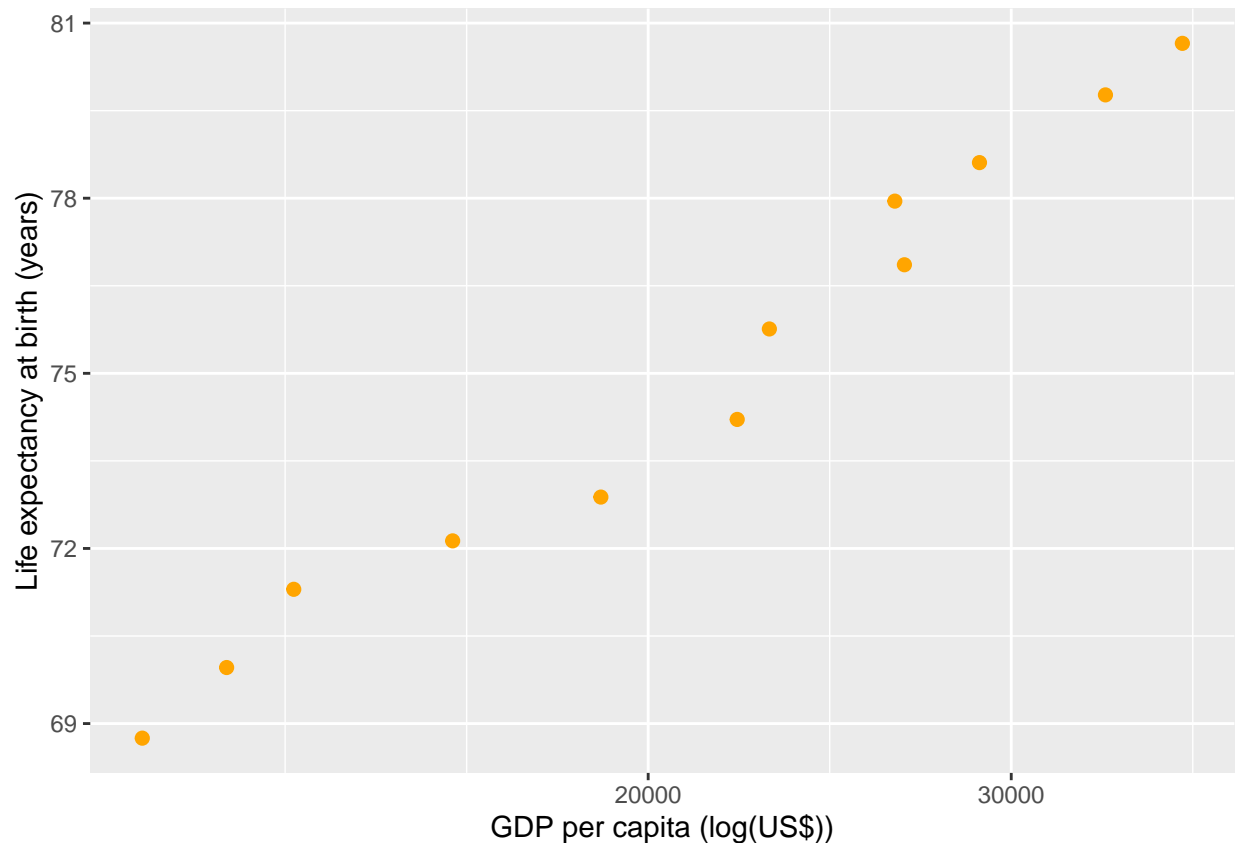
```
# A tibble: 142 x 2
# Groups:   country [142]
  country      max_gdpPercap
  <fct>         <dbl>
1 Afghanistan     978.
2 Albania         5937.
3 Algeria         6223.
4 Angola          5523.
5 Argentina       12779.
6 Australia       34435.
7 Austria         36126.
8 Bahrain         29796.
9 Bangladesh      1391.
10 Belgium        33693.
# ... with 132 more rows
```

This output allows us to determine in which year a country had its maximum GDP per capita, and what the life expectancy, population, and GDP per capita was at that time.

1.5 - Plotting

Let's investigate the relationship between life expectancy and GDP per capita in Canada. To do this, we will create a scatterplot using `ggplot2`.

```
gapminder %>%
  filter(country == "Canada") %>%
  ggplot(aes(gdpPercap, lifeExp)) +
  geom_point(colour = "orange", size = 2) +
  scale_x_log10() +
  xlab("GDP per capita (log(US$))") +
  ylab("Life expectancy at birth (years)")
```



From this plot, we can see that in Canada, life expectancy has been increasing relatively linearly with the log transform of GDP per capita between the years of 1952 and 2007.

Question 2 - Variable Exploration

To perform individual variable exploration using `dplyr` we will choose one categorical variable and one quantitative variable to explore. Let's say we want to analyze `continent` as the categorical variable and `population` as the quantitative variable.

2.1 - Categorical

The categorical variable we are interested in exploring is `continent`. To start off, let's first investigate which continents are represented in our `gapminder` dataset.

```
levels(gapminder$continent)
```

```
[1] "Africa" "Americas" "Asia" "Europe" "Oceania"
```

Categorical variable exploration is usually performed through the generation of frequency tables - let's generate one for the `continent` variable.

```
gapminder %>%
  count(continent) %>%
  knitr::kable()
```

continent	n
Africa	624
Americas	300
Asia	396
Europe	360
Oceania	24


This command prints out the number of observations for each continent in our dataset. However, by looking at the `gapminder` dataset, we see that each country contributes 12 observations (representing 12 time points) to the dataset. Suppose we want to know how many countries are in each continent for the `gapminder` dataset. To determine this, we need to remove replicates of the same country. The simplest way to achieve this is to consider continent counts at each time point.

```
gapminder %>%
  group_by(year) %>%
  count(continent)
```

```
# A tibble: 60 x 3
# Groups:   year [12]
   year continent      n
  <int> <fct>      <int>
1  1952 Africa      52
2  1952 Americas    25
3  1952 Asia       33
4  1952 Europe     30
5  1952 Oceania      2
6  1957 Africa      52
7  1957 Americas    25
8  1957 Asia       33
9  1957 Europe     30
10 1957 Oceania      2
# ... with 50 more rows
```

We see that country distribution does not change from year to year - there are the same number of countries in each of the 5 continents: Africa, the Americas, Asia, Europe, and Oceania. In our dataset, 52 countries are represented from Africa, 25 countries from the Americas, 33 countries from Asia, 30 countries from Europe, and 2 countries from Oceania. Let's view this data graphically using a bar graph - because continent counts do not change between years, it is sufficient to plot data from a single time point.

```
gapminder %>%
  group_by(year) %>%
  filter(year==2007) %>%
  ggplot(aes(continent)) +
  geom_bar(fill = "orange", alpha = 0.7) +
  ylab("Number of Countries") +
  xlab("Continent")
```



2.2 - Quantitative

Quantitative variable exploration is usually performed through the generation of 5-number summaries: min, 1st quartile, median, 3rd quartile, max. Let's generate a variation of the 5-number summary for the pop variable, including a measure of spread.

```
gapminder %>%
  summarise("min_pop (million)" = min(pop)/(10^6), "mean_pop (million)" = mean(pop)/(10^6),
            "median_pop (million)" = median(pop)/(10^6), "max_pop (million)" = max(pop)/(10^6),
            "sd_pop (million)" = sd(pop)/(10^6)) %>%
  knitr::kable()
```

min_pop (million)	mean_pop (million)	median_pop (million)	max_pop (million)	sd_pop (million)
0.060011	29.60121	7.023595	1318.683	106.1579

From the output, we can observe the minimum and maximum values for population as well as the median and mean values and the standard deviation. We observe a deviation between population mean and population median, suggesting that the distribution is skewed with greater density towards lower values ($\text{mean_pop} > \text{median_pop}$).


However, these values are difficult to interpret because they refer to the population distribution across all countries between 1952 and 2007. It may be more informative if we investigated the change in population distribution from year to year, for example.

```
gapminder %>%
  select(year, pop) %>%
  group_by(year) %>%
  summarise("min_pop (million)" = min(pop)/(10^6), "mean_pop (million)" = mean(pop)/(10^6),
            "median_pop (million)" = median(pop)/(10^6), "max_pop (million)" = max(pop)/(10^6),
            "sd_pop (million)" = sd(pop)/(10^6)) %>%
  knitr::kable()
```

year	min_pop (million)	mean_pop (million)	median_pop (million)	max_pop (million)	sd_pop (million)
1952	0.060011	16.95040	3.943953	556.2635	58.10086
1957	0.061325	18.76341	4.282942	637.4080	65.50429
1962	0.065345	20.42101	4.686039	665.7700	69.78865
1967	0.070787	22.65830	5.170176	754.5500	78.37548
1972	0.076595	25.18998	5.877997	862.0300	88.64682
1977	0.086796	27.67638	6.404037	943.4550	97.48109
1982	0.098593	30.20730	7.007320	1000.2810	105.09865
1987	0.110812	33.03857	7.774862	1084.0350	114.75618
1992	0.125911	35.99092	8.688686	1164.9700	124.50259
1997	0.145608	38.83947	9.735064	1230.0750	133.41739
2002	0.170372	41.45759	10.372919	1280.4000	140.84828
2007	0.199579	44.02122	10.517531	1318.6831	147.62140

From these summary statistics, we can see that there is a lot of variation in the `pop` variable (large standard deviation values). This is to be expected due to the large variation in population size across countries in the world. To visualize this graphically, we can plot side-by-side boxplots representing global population data at each time point.

```
gapminder %>%
  mutate(year = factor(year)) %>%
  ggplot(aes(year, pop)) +
  geom_boxplot() +
  scale_y_log10("Population (log scale)") +
  xlab("Year")
```




From this graph, we can observe a relatively consistent slow linear increase in log-transformed population over time, suggesting that between 1952 and 2007, global population has been steadily increasing.

Question 3 - Plot Exploration

3.1 - Scatterplot

The first plot type that we are going to explore is the scatterplot. Suppose we want to group the `gapminder` dataset by continent, and then plot the relationship between life expectancy at birth and GDP per capita.

```
gapminder %>%
  group_by(continent) %>%
  ggplot(aes(x = gdpPercap, y = lifeExp, colour = continent)) +
  geom_point(alpha = 0.7) +
  scale_x_log10() +
  ylab("Life expectancy at birth (year)") +
  xlab("GDP per capita (log(US$))")
```





From the scatterplot, we can see that life expectancy and GDP per capita are roughly positively correlated, following a linear relationship. Not only does life expectancy tend to increase with GDP per capita, we can also observe that countries in Europe and Oceania tend to have higher life expectancy and GDP whereas countries in Africa tend to have lower life expectancy and GDP. Countries in the Americas and Asia seem to have more varied distributions of life expectancy and GDP with more discrepancy between countries.

3.2 - Time Series

The second plot type we are going to explore is a time series. Suppose we want to investigate how population has changed over time between 1952 and 2007 in Canada, France, and Australia. We can visualize this through making a time series plot and fitting a trendline.

```
gapminder %>%
  filter(country == "Canada" | country == "France" | country == "Australia") %>%
  ggplot(aes(x = year, y = pop, col = country)) +
  geom_point(size = 2) +
  geom_line(size = 1) +
  xlab("Year") +
  ylab("Population") +
  theme(legend.title=element_blank())
```



From observing the trendlines, we can observe that all countries have experienced roughly linear increases in population over time. Canada and France seem to have similar rates of population increase, while the rate of increase in population is slower in Australia. Between 1952 and 2007, Canada and France have experienced an increase in population of approximately 20 million people, while Australia's population has increased by approximately 10 million people.

Bonus - Recycling

Evaluation of the following command:

```
filter(gapminder, country == c("Rwanda", "Afghanistan"))
```

```
# A tibble: 12 x 6
  country    continent  year lifeExp    pop gdpPercap
  <fct>      <fct>    <int>  <dbl>   <int>   <dbl>
1 Afghanistan Asia      1957   30.3  9240934    821.
2 Afghanistan Asia      1967   34.0 11537966    836.
3 Afghanistan Asia      1977   38.4 14880372    786.
4 Afghanistan Asia      1987   40.8 13867957    852.
5 Afghanistan Asia      1997   41.8 22227415    635.
6 Afghanistan Asia      2007   43.8 31889923    975.
7 Rwanda     Africa    1952   40    2534927    493.
8 Rwanda     Africa    1962   43    3051242    597.
```

9	Rwanda	Africa	1972	44.6	3992121	591.
10	Rwanda	Africa	1982	46.2	5507565	882.
11	Rwanda	Africa	1992	23.6	7290203	737.
12	Rwanda	Africa	2002	43.4	7852401	786.

The analyst's goal was to obtain a subset of the `gapminder` dataset, containing data from Rwanda and Afghanistan. However, the output of the above command only returns 12 observations when there should have been 24 (12 observations for each country representing the 12 time points between 1952 and 2007). Each year should have both a Rwanda observation and an Afghanistan observation; however, in our subset, each year is only represented by one country. This suggests that the subsetting data is incomplete.

To correctly obtain all data from Rwanda and Afghanistan, we can use the logical “or” operator (denoted as “|” in `dplyr`). The code below will return a dataframe that contains all observations from either Rwanda or Afghanistan.

```
filter(gapminder, country == "Rwanda" | country == "Afghanistan")
```

```
# A tibble: 24 x 6
  country    continent  year lifeExp      pop gdpPercap
  <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
1 Afghanistan Asia      1952   28.8  8425333    779.
2 Afghanistan Asia      1957   30.3  9240934    821.
3 Afghanistan Asia      1962   32.0 10267083    853.
4 Afghanistan Asia      1967   34.0 11537966    836.
5 Afghanistan Asia      1972   36.1 13079460    740.
6 Afghanistan Asia      1977   38.4 14880372    786.
7 Afghanistan Asia      1982   39.9 12881816    978.
8 Afghanistan Asia      1987   40.8 13867957    852.
9 Afghanistan Asia      1992   41.7 16317921    649.
10 Afghanistan Asia      1997   41.8 22227415    635.
# ... with 14 more rows
```

From this output, we see that we are left with a dataframe of the correct dimensions (24 observations) with all data from Rwanda and Afghanistan between 1952 and 2007.